

Data Wrangling with dplyr

Abhinav Anand, IIMB

2018/06/08

Setup

The package suite `tidyverse()` which includes the package `dplyr` needs to be included. Also, the package `gapminder` needs to be installed prior to running the commands below. For installing `dplyr`, type `install.packages("dplyr")` or equivalently for `tidyverse`, type `install.packages("tidyverse")`. To install `gapminder`, type `install.packages("gapminder")` in the RStudio console.

```
(data_gapminder <- gapminder::gapminder)
```

```
## # A tibble: 1,704 x 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
## 7 Afghanistan Asia      1982   39.9 12881816    978.
## 8 Afghanistan Asia      1987   40.8 13867957    852.
## 9 Afghanistan Asia      1992   41.7 16317921    649.
## 10 Afghanistan Asia      1997   41.8 22227415    635.
## # ... with 1,694 more rows
```

```
#what is the format: wide or long?
```

```
#what is a tibble?
```

Notes

1. **Data Types:** R has many in-built data types. Examples:
 - i. **fct:** “factor”: categorical data which can assume finite levels, say A, B, C etc.
 - ii. **dbl:** “double”: real numbers, say 3.671, 4.00, 10.122482929 etc.
 - iii. **int:** “integer”: integers, say 3, 10, -9, 0 etc.
 - iv. **chr:** “character”: say, “FMC”, “Term 4” etc.
 - v. **lgl:** “logical”: $\{0, 1\} \equiv \{T, F\}$
 - vi. **date**, and many more
2. **Tibbles:** Tibbles are essentially data frames, but slightly altered to work better in tidyverse. (Compare `head(data_frame_name)` versus `tibble_name`.)

dplyr(): The Main Verbs

1. **filter():** Extract rows
2. **select():** Extract columns
3. **arrange():** Order rows
4. **mutate():** Create new columns (= variables)
5. **summarise():** Compute summary statistics

The syntax for all five verbs is similar. The first argument is the data frame, followed by the action to be performed using the variable name.

Filter

Let us observe the state of the world in 1952 and the contrast and compare with that in 2007.

```
(data_1952 <- data_gapminder %>%  
  dplyr::filter(year == 1952) #extract the rows for year 1952  
 ) #note == as opposed to =
```

```
## # A tibble: 142 x 6
```

```
##   country      continent  year lifeExp      pop gdpPercap
```

```
##      <fct>      <fct>      <int>    <dbl>      <int>      <dbl>
##  1 Afghanistan Asia        1952     28.8  8425333     779.
##  2 Albania      Europe    1952     55.2  1282697    1601.
##  3 Algeria      Africa    1952     43.1  9279525    2449.
##  4 Angola       Africa    1952     30.0  4232095    3521.
##  5 Argentina    Americas  1952     62.5  17876956    5911.
##  6 Australia    Oceania   1952     69.1  8691212   10040.
##  7 Austria      Europe    1952     66.8  6927772    6137.
##  8 Bahrain      Asia      1952     50.9   120447    9867.
##  9 Bangladesh   Asia      1952     37.5  46886859     684.
## 10 Belgium      Europe    1952     68    8730405    8343.
## # ... with 132 more rows
```

```
(data_2007 <- data_gapminder %>%
  dplyr::filter(year == 2007) #extract the rows for year 2007
)
```

```
## # A tibble: 142 x 6
##   country      continent year lifeExp      pop gdpPercap
##   <fct>        <fct>      <int>    <dbl>      <int>      <dbl>
##  1 Afghanistan Asia        2007     43.8  31889923     975.
##  2 Albania      Europe    2007     76.4   3600523    5937.
##  3 Algeria      Africa    2007     72.3  33333216    6223.
##  4 Angola       Africa    2007     42.7  12420476    4797.
##  5 Argentina    Americas  2007     75.3  40301927   12779.
##  6 Australia    Oceania   2007     81.2  20434176   34435.
##  7 Austria      Europe    2007     79.8   8199783   36126.
##  8 Bahrain      Asia      2007     75.6    708573   29796.
##  9 Bangladesh   Asia      2007     64.1 150448339    1391.
## 10 Belgium      Europe    2007     79.4  10392226   33693.
## # ... with 132 more rows
```

Select

Let us also focus on two variables—GDP/capita and life expectancy. We extract both for years 1952 and 2007.

```
(data_1952_gdppc <- data_1952 %>%  
  dplyr::select(country, year, gdpPercap)  
)
```

```
## # A tibble: 142 x 3  
##   country      year gdpPercap  
##   <fct>      <int>    <dbl>  
## 1 Afghanistan 1952      779.  
## 2 Albania      1952     1601.  
## 3 Algeria      1952     2449.  
## 4 Angola       1952     3521.  
## 5 Argentina    1952     5911.  
## 6 Australia    1952    10040.  
## 7 Austria      1952     6137.  
## 8 Bahrain      1952     9867.  
## 9 Bangladesh   1952      684.  
## 10 Belgium     1952     8343.  
## # ... with 132 more rows
```

```
(data_2007_gdppc <- data_2007 %>%  
  dplyr::select(country, year, gdpPercap)  
)
```

```
## # A tibble: 142 x 3  
##   country      year gdpPercap  
##   <fct>      <int>    <dbl>  
## 1 Afghanistan 2007      975.  
## 2 Albania      2007     5937.  
## 3 Algeria      2007     6223.  
## 4 Angola       2007     4797.  
## 5 Argentina    2007    12779.
```

```
## 6 Australia      2007      34435.
## 7 Austria        2007      36126.
## 8 Bahrain        2007      29796.
## 9 Bangladesh     2007       1391.
## 10 Belgium       2007      33693.
## # ... with 132 more rows
```

```
(data_1952_life_exp <- data_1952 %>%
  dplyr::select(country, year, lifeExp)
)
```

```
## # A tibble: 142 x 3
##   country      year lifeExp
##   <fct>      <int>   <dbl>
## 1 Afghanistan  1952     28.8
## 2 Albania      1952     55.2
## 3 Algeria      1952     43.1
## 4 Angola       1952     30.0
## 5 Argentina    1952     62.5
## 6 Australia    1952     69.1
## 7 Austria      1952     66.8
## 8 Bahrain      1952     50.9
## 9 Bangladesh   1952     37.5
## 10 Belgium     1952     68
## # ... with 132 more rows
```

```
(data_2007_life_exp <- data_2007 %>%
  dplyr::select(country, year, lifeExp)
)
```

```
## # A tibble: 142 x 3
##   country      year lifeExp
##   <fct>      <int>   <dbl>
## 1 Afghanistan  2007     43.8
## 2 Albania      2007     76.4
## 3 Algeria      2007     72.3
```

```
## 4 Angola      2007    42.7
## 5 Argentina   2007    75.3
## 6 Australia   2007    81.2
## 7 Austria     2007    79.8
## 8 Bahrain     2007    75.6
## 9 Bangladesh  2007    64.1
## 10 Belgium    2007    79.4
## # ... with 132 more rows
```

```
# = dplyr::select(-c(continent, pop, gdpPercap))
```

`dplyr::rename()` is a wrapper function for `select()` which renames the variable in consideration and keeps all other variables intact.

Arrange

Usage of `arrange()` orders (from first to last) entries on the basis of a variable.

Question: Is the set of richest countries the same in 1952 and 2007?

```
(data_1952_rich <- data_1952_gdppc %>%
  dplyr::arrange(desc(gdpPercap)) #note the use of desc()
)
```

```
## # A tibble: 142 x 3
##   country      year gdpPercap
##   <fct>      <int>    <dbl>
## 1 Kuwait      1952  108382.
## 2 Switzerland 1952   14734.
## 3 United States 1952   13990.
## 4 Canada      1952   11367.
## 5 New Zealand  1952   10557.
## 6 Norway      1952   10095.
## 7 Australia    1952   10040.
## 8 United Kingdom 1952    9980.
```

```
## 9 Bahrain          1952      9867.
## 10 Denmark         1952      9692.
## # ... with 132 more rows
```

```
(data_2007_rich <- data_2007_gdppc %>%
  dplyr::arrange(desc(gdpPercap)) #note the use of desc()
)
```

```
## # A tibble: 142 x 3
##   country          year gdpPercap
##   <fct>          <int>     <dbl>
## 1 Norway          2007     49357.
## 2 Kuwait          2007     47307.
## 3 Singapore       2007     47143.
## 4 United States   2007     42952.
## 5 Ireland         2007     40676.
## 6 Hong Kong, China 2007     39725.
## 7 Switzerland     2007     37506.
## 8 Netherlands     2007     36798.
## 9 Canada          2007     36319.
## 10 Iceland        2007     36181.
## # ... with 132 more rows
```

Which countries display the highest life expectancy pre-2000?

```
(data_life_pre00 <- data_gapminder %>%
  dplyr::select(country, year, lifeExp) %>%
  dplyr::filter(year <= 2000) %>%
  dplyr::arrange(desc(lifeExp))
)
```

```
## # A tibble: 1,420 x 3
##   country          year lifeExp
##   <fct>          <int>     <dbl>
## 1 Japan          1997      80.7
```

```
## 2 Hong Kong, China 1997 80
## 3 Sweden          1997 79.4
## 4 Switzerland     1997 79.4
## 5 Japan            1992 79.4
## 6 Iceland         1997 79.0
## 7 Australia       1997 78.8
## 8 Italy            1997 78.8
## 9 Iceland         1992 78.8
## 10 Spain           1997 78.8
## # ... with 1,410 more rows
```

Post-2000?

```
(data_life_post00 <- data_gapminder %>%
  dplyr::select(country, year, lifeExp) %>%
  dplyr::filter(year > 2000) %>%
  dplyr::arrange(desc(lifeExp))
)
```

```
## # A tibble: 284 x 3
##   country      year lifeExp
##   <fct>      <int>   <dbl>
## 1 Japan      2007    82.6
## 2 Hong Kong, China 2007    82.2
## 3 Japan      2002    82
## 4 Iceland    2007    81.8
## 5 Switzerland 2007    81.7
## 6 Hong Kong, China 2002    81.5
## 7 Australia   2007    81.2
## 8 Spain       2007    80.9
## 9 Sweden      2007    80.9
## 10 Israel     2007    80.7
## # ... with 274 more rows
```


Mutate

Let's define a new variable called "total GDP" which is the product of the GDP/capita and the total population. To compute it and include it in the list of variables we can use the verb `dplyr::mutate()`.

```
(data_GDP_tot <- data_gapminder %>%  
  dplyr::mutate(GDP_total = pop*gdpPercap/10^9) #in USD billions  
)
```



```
## # A tibble: 1,704 x 7  
##   country      continent  year lifeExp      pop gdpPercap GDP_total  
##   <fct>         <fct>    <int>  <dbl>    <int>    <dbl>    <dbl>  
## 1 Afghanistan Asia      1952   28.8  8425333    779.     6.57  
## 2 Afghanistan Asia      1957   30.3  9240934    821.     7.59  
## 3 Afghanistan Asia      1962   32.0 10267083    853.     8.76  
## 4 Afghanistan Asia      1967   34.0 11537966    836.     9.65  
## 5 Afghanistan Asia      1972   36.1 13079460    740.     9.68  
## 6 Afghanistan Asia      1977   38.4 14880372    786.    11.7  
## 7 Afghanistan Asia      1982   39.9 12881816    978.    12.6  
## 8 Afghanistan Asia      1987   40.8 13867957    852.    11.8  
## 9 Afghanistan Asia      1992   41.7 16317921    649.    10.6  
## 10 Afghanistan Asia     1997   41.8 22227415    635.    14.1  
## # ... with 1,694 more rows
```

In general, for `mutate()` to work well, the function must take a vector of values as input and return a vector with the same number of values as output. A short list of functions that can be used with `mutate()` are:

1. Arithmetic operators: `+`, `-`, `*`, `/`, `^`
2. Logs: `log()`, `log2()`, `log10()`
3. Cumulative aggregates: `cumsum()`, `cumprod()`, `cummin()`, `cummax()`, `cummean()` etc.

and many more.

Question: Which countries have the highest total GDP in 1952 and 2007?

```
(data_GDP_tot %>%  
  dplyr::filter(year == 1952) %>%  
  dplyr::arrange(desc(GDP_total))  
)
```

```
## # A tibble: 142 x 7  
##   country      continent  year lifeExp      pop gdpPercap GDP_total  
##   <fct>        <fct>    <int> <dbl>    <int>    <dbl>    <dbl>  
## 1 United States Americas  1952  68.4 157553000 13990.  2204.  
## 2 United Kingdom Europe    1952  69.2  50430000  9980.   503.  
## 3 Germany      Europe    1952  67.5  69145952  7144.   494.  
## 4 France        Europe    1952  67.4  42459667  7030.   298.  
## 5 Japan         Asia     1952  63.0  86459025  3217.   278.  
## 6 Italy          Europe    1952  65.9  47666000  4931.   235.  
## 7 China         Asia     1952   44  556263527  400.    223.  
## 8 India         Asia     1952  37.4 372000000  547.    203.  
## 9 Canada        Americas  1952  68.8  14785584 11367.   168.  
## 10 Brazil       Americas  1952  50.9  56602560  2109.   119.  
## # ... with 132 more rows
```

```
(data_GDP_tot %>%  
  dplyr::filter(year == 2007) %>%  
  dplyr::arrange(desc(GDP_total))  
)
```

```
## # A tibble: 142 x 7  
##   country      continent  year lifeExp      pop gdpPercap GDP_total  
##   <fct>        <fct>    <int> <dbl>    <int>    <dbl>    <dbl>  
## 1 United States Americas  2007  78.2  301139947 42952. 12934.  
## 2 China         Asia     2007  73.0 1318683096  4959.  6540.  
## 3 Japan         Asia     2007  82.6  127467972 31656.  4035.  
## 4 India         Asia     2007  64.7 1110396331  2452.  2723.  
## 5 Germany      Europe    2007  79.4  82400996  32170. 2651.
```

```
## 6 United Kingdom Europe 2007 79.4 60776238 33203. 2018.
## 7 France Europe 2007 80.7 61083916 30470. 1861.
## 8 Brazil Americas 2007 72.4 190010647 9066. 1723.
## 9 Italy Europe 2007 80.5 58147733 28570. 1661.
## 10 Mexico Americas 2007 76.2 108700891 11978. 1302.
## # ... with 132 more rows
```

Question: Which are the five smallest economies in 2007 (by total GDP)?

```
(data_GDP_tot %>%
  dplyr::filter(year == 2007) %>%
  dplyr::arrange(GDP_total) %>%
  dplyr::filter(rank(GDP_total) <= 5)
)
```

```
## # A tibble: 5 x 7
##   country          continent year lifeExp   pop gdpPercap GDP_total
##   <fct>           <fct>    <int> <dbl> <int>    <dbl>    <dbl>
## 1 Sao Tome and Principe Africa  2007  65.5 2.00e5  1598.    0.319
## 2 Comoros         Africa  2007  65.2 7.11e5   986.    0.701
## 3 Guinea-Bissau   Africa  2007  46.4 1.47e6   579.    0.853
## 4 Djibouti        Africa  2007  54.8 4.96e5  2082.    1.03
## 5 Gambia          Africa  2007  59.4 1.69e6   753.    1.27
```

Summarise

The function `summarise()` (or equivalently `summarize()`) can be used to compute summary statistics. Here is an example, where we summarize the variable life expectancy for the continent Europe.

```
(life_exp_summ_eur <- data_gapminder %>%
  dplyr::filter(continent == "Europe") %>%
  dplyr::summarise(average = mean(lifeExp),
    med = median(lifeExp),
```

```

        std = sd(lifeExp),
        variance = var(lifeExp),
        iqr = IQR(lifeExp)
      )
    )
  )
)

```

```

## # A tibble: 1 x 5
##   average med   std variance   iqr
##   <dbl> <dbl> <dbl>     <dbl> <dbl>
## 1   71.9  72.2  5.43     29.5  5.88

```

Grouped Summaries

Question: What are continent-wise summary statistics?

```

(summ_life_exp <- data_gapminder %>%
  dplyr::group_by(continent) %>%
  dplyr::summarise(average = mean(lifeExp),
                    med = median(lifeExp),
                    std = sd(lifeExp),
                    variance = var(lifeExp),
                    iqr = IQR(lifeExp)
  )
)


```

```

## # A tibble: 5 x 6
##   continent average   med   std variance   iqr
##   <fct>         <dbl> <dbl> <dbl>     <dbl> <dbl>
## 1 Africa       48.9  47.8  9.15     83.7  12.0
## 2 Americas     64.7  67.0  9.35     87.3  13.3
## 3 Asia         60.1  61.8 11.9     141.   18.1
## 4 Europe       71.9  72.2  5.43     29.5   5.88
## 5 Oceania      74.3  73.7  3.80     14.4   6.35

```

```
ggplot(data = summ_life_exp,
       aes(x = continent, y = average)) +
  geom_bar(stat = "identity")
```

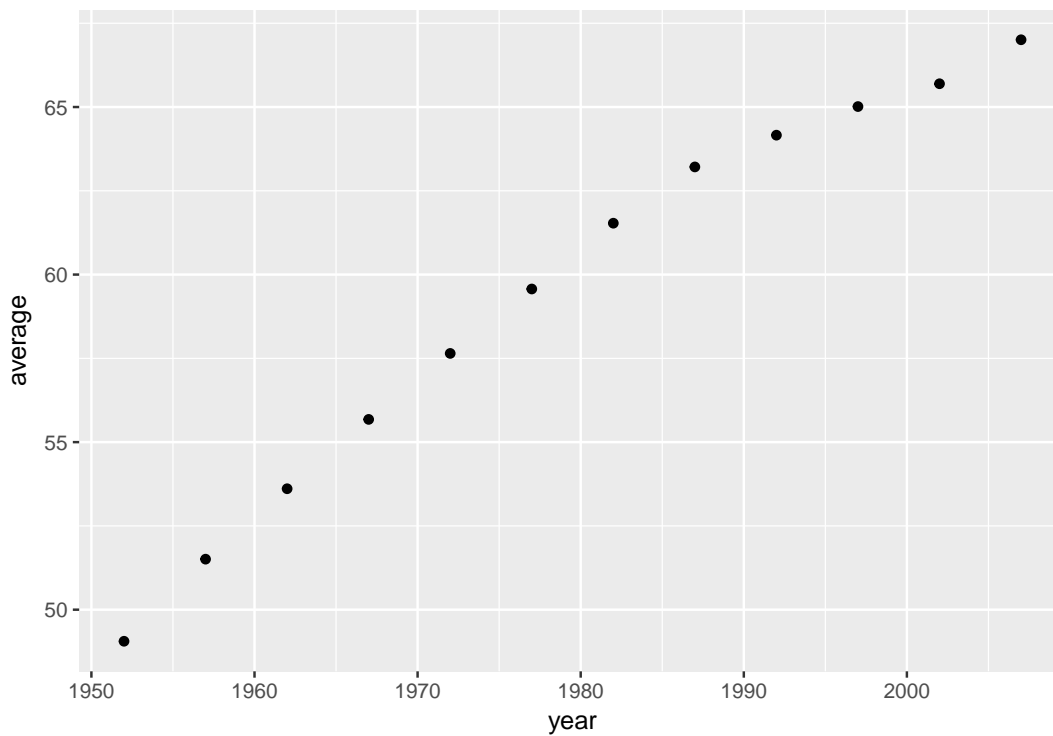


```
(summ_year_life_exp <- data_gapminder %>%
  dplyr::group_by(year) %>%
  dplyr::summarise(average = mean(lifeExp),
                  med = median(lifeExp),
                  std = sd(lifeExp),
                  variance = var(lifeExp),
                  iqr = IQR(lifeExp)
                )
)
```

```
## # A tibble: 12 x 6
##   year average med std variance iqr
##   <int> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 1952 49.1 45.1 12.2 149. 20.7
```

##	2	1957	51.5	48.4	12.2	150.	21.8
##	3	1962	53.6	50.9	12.1	146.	21.8
##	4	1967	55.7	53.8	11.7	137.	21.4
##	5	1972	57.6	56.5	11.4	130.	20.7
##	6	1977	59.6	59.7	11.2	126.	19.9
##	7	1982	61.5	62.4	10.8	116.	18.0
##	8	1987	63.2	65.8	10.6	111.	16.9
##	9	1992	64.2	67.7	11.2	126.	16.5
##	10	1997	65.0	69.4	11.6	134.	18.5
##	11	2002	65.7	70.8	12.3	151.	19.9
##	12	2007	67.0	71.9	12.1	146.	19.3

```
ggplot(data = summ_year_life_exp,
       aes(x = year, y = average)) +
  geom_point()
```



One can also perform grouped mutates and grouped filters.