

Data Wrangling with dplyr

Abhinav Anand

For data wrangling, the package `dplyr` is employed. For downloading, type `install.packages("dplyr")` or equivalently type `install.packages("tidyverse")`

```
(data_gapminder <- gapminder::gapminder)
```

```
## # A tibble: 1,704 x 6
##   country      continent year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>   <dbl>   <int>   <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
## 7 Afghanistan Asia      1982   39.9 12881816    978.
## 8 Afghanistan Asia      1987   40.8 13867957    852.
## 9 Afghanistan Asia      1992   41.7 16317921    649.
## 10 Afghanistan Asia      1997   41.8 22227415    635.
## # ... with 1,694 more rows
```

```
#what is the format: wide or long?
#what is a tibble?
```

Notes

1. **Data Types:** R has many in-built data types. Examples:
 - i. `fct`: “factor”: categorical data which can assume finite levels, say A, B, C etc.
 - ii. `dbl`: “double”: real numbers, say 3.671, 4.00, 10.122482929 etc.
 - iii. `int`: “integer”: integers, say 3, 10, -9, 0 etc.
 - iv. `chr`: “character”: say, “FMC”, “Term 4” etc.

- v. `lgl`: “logical”: $\{0, 1\} \equiv \{T, F\}$
- vi. `date`, and many more
- 2. **Tibbles**: Tibbles are essentially data frames, but slightly altered to work better in tidyverse. (Compare `head(data_frame_name)` versus `tibble_name`.)

dplyr(): The Main Verbs

1. `filter()`: Extract rows
2. `select()`: Extract columns
3. `arrange()`: Order rows
4. `mutate()`: Create new columns (= variables)
5. `summarise()`: Compute summary statistics

The syntax for all five verbs is similar. The first argument is the data frame, followed by the action to be performed using the variable name.

Filter

Let us observe the state of the world in 1952 and the contrast and compare with that in 2007.

```
(data_1952 <- data_gapminder %>%
  dplyr::filter(year == 1952) #extract the rows for year 1952
) #note == as opposed to =
```

```
## # A tibble: 142 x 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>         <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Albania      Europe    1952   55.2  1282697   1601.
## 3 Algeria      Africa    1952   43.1  9279525   2449.
## 4 Angola       Africa    1952   30.0  4232095   3521.
## 5 Argentina    Americas  1952   62.5  17876956  5911.
## 6 Australia    Oceania   1952   69.1  8691212  10040.
## 7 Austria      Europe    1952   66.8  6927772   6137.
```

```
## 8 Bahrain      Asia      1952    50.9  120447    9867.
## 9 Bangladesh   Asia      1952    37.5 46886859    684.
## 10 Belgium     Europe    1952    68    8730405    8343.
## # ... with 132 more rows
```

```
(data_2007 <- data_gapminder %>%
  dplyr::filter(year == 2007) #extract the rows for year 2007
)
```

```
## # A tibble: 142 x 6
##   country      continent year lifeExp      pop gdpPercap
##   <fct>         <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      2007   43.8 31889923    975.
## 2 Albania      Europe    2007   76.4  3600523   5937.
## 3 Algeria      Africa    2007   72.3 33333216   6223.
## 4 Angola       Africa    2007   42.7 12420476   4797.
## 5 Argentina    Americas  2007   75.3 40301927  12779.
## 6 Australia    Oceania   2007   81.2 20434176  34435.
## 7 Austria      Europe    2007   79.8  8199783   36126.
## 8 Bahrain      Asia      2007   75.6   708573   29796.
## 9 Bangladesh   Asia      2007   64.1 150448339  1391.
## 10 Belgium     Europe    2007   79.4 10392226   33693.
## # ... with 132 more rows
```

Select

Let us also focus on two variables—GDP/capita and life expectancy. We extract both for years 1952 and 2007.

```
(data_1952_gdppc <- data_1952 %>%
  dplyr::select(country, year, gdpPercap)
)
```

```
## # A tibble: 142 x 3
##   country      year gdpPercap
##   <fct>         <int>    <dbl>
## 1 Afghanistan 1952    779.
```

```
## 2 Albania      1952      1601.
## 3 Algeria      1952      2449.
## 4 Angola       1952      3521.
## 5 Argentina    1952      5911.
## 6 Australia    1952     10040.
## 7 Austria      1952      6137.
## 8 Bahrain      1952      9867.
## 9 Bangladesh   1952       684.
## 10 Belgium     1952     8343.
## # ... with 132 more rows
```

```
(data_2007_gdppc <- data_2007 %>%
  dplyr::select(country, year, gdpPercap)
)
```

```
## # A tibble: 142 x 3
##   country      year gdpPercap
##   <fct>      <int>    <dbl>
## 1 Afghanistan 2007      975.
## 2 Albania     2007     5937.
## 3 Algeria     2007     6223.
## 4 Angola      2007     4797.
## 5 Argentina   2007    12779.
## 6 Australia   2007    34435.
## 7 Austria     2007    36126.
## 8 Bahrain     2007    29796.
## 9 Bangladesh  2007     1391.
## 10 Belgium    2007    33693.
## # ... with 132 more rows
```

```
(data_1952_life_exp <- data_1952 %>%
  dplyr::select(country, year, lifeExp)
)
```

```
## # A tibble: 142 x 3
##   country      year lifeExp
##   <fct>      <int>    <dbl>
```

```
## 1 Afghanistan 1952 28.8
## 2 Albania      1952 55.2
## 3 Algeria      1952 43.1
## 4 Angola       1952 30.0
## 5 Argentina    1952 62.5
## 6 Australia    1952 69.1
## 7 Austria      1952 66.8
## 8 Bahrain      1952 50.9
## 9 Bangladesh   1952 37.5
## 10 Belgium     1952 68
## # ... with 132 more rows
```

```
(data_2007_life_exp <- data_2007 %>%
  dplyr::select(country, year, lifeExp)
)
```

```
## # A tibble: 142 x 3
##   country      year lifeExp
##   <fct>      <int>   <dbl>
## 1 Afghanistan 2007    43.8
## 2 Albania     2007    76.4
## 3 Algeria     2007    72.3
## 4 Angola      2007    42.7
## 5 Argentina   2007    75.3
## 6 Australia   2007    81.2
## 7 Austria     2007    79.8
## 8 Bahrain     2007    75.6
## 9 Bangladesh  2007    64.1
## 10 Belgium    2007    79.4
## # ... with 132 more rows
```

```
# = dplyr::select(-c(continent, pop, gdpPercap))
```

`dplyr::rename()` is a wrapper function for `select()` which renames the variable in consideration and keeps all other variables intact.

Arrange

Usage of `arrange()` orders (from first to last) entries on the basis of a variable.

Question: Is the set of richest countries the same in 1952 and 2007?

```
(data_1952_rich <- data_1952_gdppc %>%  
  dplyr::arrange(desc(gdpPercap)) #note the use of desc()  
)
```

```
## # A tibble: 142 x 3  
##   country      year gdpPercap  
##   <fct>      <int>    <dbl>  
## 1 Kuwait      1952    108382.  
## 2 Switzerland 1952     14734.  
## 3 United States 1952     13990.  
## 4 Canada       1952     11367.  
## 5 New Zealand  1952     10557.  
## 6 Norway       1952      10095.  
## 7 Australia    1952      10040.  
## 8 United Kingdom 1952       9980.  
## 9 Bahrain      1952       9867.  
## 10 Denmark     1952       9692.  
## # ... with 132 more rows
```

```
(data_2007_rich <- data_2007_gdppc %>%  
  dplyr::arrange(desc(gdpPercap)) #note the use of desc()  
)
```

```
## # A tibble: 142 x 3  
##   country      year gdpPercap  
##   <fct>      <int>    <dbl>  
## 1 Norway      2007    49357.  
## 2 Kuwait      2007    47307.  
## 3 Singapore   2007    47143.  
## 4 United States 2007    42952.
```

```
## 5 Ireland          2007    40676.
## 6 Hong Kong, China 2007    39725.
## 7 Switzerland      2007    37506.
## 8 Netherlands      2007    36798.
## 9 Canada           2007    36319.
## 10 Iceland         2007    36181.
## # ... with 132 more rows
```

Which countries display the highest life expectancy pre-2000?

```
(data_life_pre00 <- data_gapminder %>%
  dplyr::select(country, year, lifeExp) %>%
  dplyr::filter(year <= 2000) %>%
  dplyr::arrange(desc(lifeExp))
)
```

```
## # A tibble: 1,420 x 3
##   country      year lifeExp
##   <fct>      <int>   <dbl>
## 1 Japan      1997    80.7
## 2 Hong Kong, China 1997    80
## 3 Sweden     1997    79.4
## 4 Switzerland 1997    79.4
## 5 Japan      1992    79.4
## 6 Iceland    1997    79.0
## 7 Australia  1997    78.8
## 8 Italy       1997    78.8
## 9 Iceland    1992    78.8
## 10 Spain     1997    78.8
## # ... with 1,410 more rows
```

Post-2000?

```
(data_life_post00 <- data_gapminder %>%
  dplyr::select(country, year, lifeExp) %>%
  dplyr::filter(year > 2000) %>%
  dplyr::arrange(desc(lifeExp))
)
```

```
## # A tibble: 284 x 3
##   country      year lifeExp
##   <fct>      <int>   <dbl>
## 1 Japan      2007    82.6
## 2 Hong Kong, China 2007    82.2
## 3 Japan      2002     82
## 4 Iceland    2007    81.8
## 5 Switzerland 2007    81.7
## 6 Hong Kong, China 2002    81.5
## 7 Australia   2007    81.2
## 8 Spain       2007    80.9
## 9 Sweden      2007    80.9
## 10 Israel     2007    80.7
## # ... with 274 more rows
```

Mutate

Let's define a new variable called "total GDP" which is the product of the GDP/capita and the total population. To compute it and include it in the list of variables we can use the verb `dplyr::mutate()`.

```
(data_GDP_tot <- data_gapminder %>%
  dplyr::mutate(GDP_total = pop*gdpPercap/10^9) #in USD billions
)
```

```
## # A tibble: 1,704 x 7
##   country      continent year lifeExp      pop gdpPercap GDP_total
##   <fct>      <fct>      <int>   <dbl>   <int>    <dbl>    <dbl>
## 1 Afghanistan Asia      1952    28.8  8425333    779.    6.57
```



```
## 2 Afghanistan Asia      1957    30.3  9240934      821.      7.59
## 3 Afghanistan Asia      1962    32.0 10267083      853.      8.76
## 4 Afghanistan Asia      1967    34.0 11537966      836.      9.65
## 5 Afghanistan Asia      1972    36.1 13079460      740.      9.68
## 6 Afghanistan Asia      1977    38.4 14880372      786.     11.7
## 7 Afghanistan Asia      1982    39.9 12881816      978.     12.6
## 8 Afghanistan Asia      1987    40.8 13867957      852.     11.8
## 9 Afghanistan Asia      1992    41.7 16317921      649.     10.6
## 10 Afghanistan Asia     1997    41.8 22227415      635.     14.1
## # ... with 1,694 more rows
```

In general, for `mutate()` to work well, the function must take a vector of values as input and return a vector with the same number of values as output. A short list of functions that can be used with `mutate()` are:

1. Arithmetic operators: `+`, `-`, `*`, `/`, `^`
2. Logs: `log()`, `log2()`, `log10()`
3. Cumulative aggregates: `cumsum()`, `cumprod()`, `cummin()`, `cummax()`, `cummean()` etc.

and many more.

Question: Which countries have the highest total GDP in 1952 and 2007?

```
(data_GDP_tot %>%
  dplyr::filter(year == 1952) %>%
  dplyr::arrange(desc(GDP_total))
)
```

```
## # A tibble: 142 x 7
##   country      continent  year lifeExp      pop gdpPercap GDP_total
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>    <dbl>
## 1 United States Americas   1952   68.4 157553000 13990.   2204.
## 2 United Kingdom Europe     1952   69.2  50430000  9980.    503.
## 3 Germany      Europe     1952   67.5  69145952  7144.    494.
## 4 France       Europe     1952   67.4  42459667  7030.    298.
## 5 Japan        Asia      1952   63.0  86459025  3217.    278.
```

```
## 6 Italy          Europe    1952    65.9  47666000    4931.    235.
## 7 China          Asia      1952     44   556263527    400.    223.
## 8 India          Asia      1952    37.4  372000000    547.    203.
## 9 Canada         Americas  1952    68.8  14785584    11367.   168.
## 10 Brazil        Americas  1952    50.9   56602560    2109.   119.
## # ... with 132 more rows
```

```
(data_GDP_tot %>%
  dplyr::filter(year == 2007) %>%
  dplyr::arrange(desc(GDP_total))
)
```

```
## # A tibble: 142 x 7
##   country      continent  year lifeExp      pop gdpPercap GDP_total
##   <fct>        <fct>    <int> <dbl>    <int>    <dbl>    <dbl>
## 1 United States Americas   2007   78.2  301139947  42952.  12934.
## 2 China        Asia     2007   73.0  1318683096  4959.   6540.
## 3 Japan        Asia     2007   82.6  127467972  31656.  4035.
## 4 India        Asia     2007   64.7  1110396331  2452.   2723.
## 5 Germany      Europe   2007   79.4   82400996  32170.  2651.
## 6 United Kingdom Europe   2007   79.4   60776238  33203.  2018.
## 7 France       Europe   2007   80.7   61083916  30470.  1861.
## 8 Brazil       Americas  2007   72.4  190010647   9066.  1723.
## 9 Italy        Europe   2007   80.5   58147733  28570.  1661.
## 10 Mexico      Americas  2007   76.2  108700891  11978.  1302.
## # ... with 132 more rows
```

Question: Which are the five smallest economies in 2007 (by total GDP)?

```
(data_GDP_tot %>%
  dplyr::filter(year == 2007) %>%
  dplyr::arrange(GDP_total) %>%
  dplyr::filter(rank(GDP_total) <= 5)
)
```

```
## # A tibble: 5 x 7
```

```
##   country          continent  year lifeExp    pop gdpPercap GDP_total
##   <fct>            <fct>    <int>  <dbl>  <int>    <dbl>    <dbl>
## 1 Sao Tome and Principe Africa    2007   65.5 2.00e5    1598.    0.319
## 2 Comoros           Africa    2007   65.2 7.11e5     986.    0.701
## 3 Guinea-Bissau      Africa    2007   46.4 1.47e6     579.    0.853
## 4 Djibouti           Africa    2007   54.8 4.96e5    2082.    1.03
## 5 Gambia             Africa    2007   59.4 1.69e6     753.    1.27
```

Summarise

The function `summarise()` (or equivalently `summarize()`) can be used to compute summary statistics. Here is an example, where we summarize the variable life expectancy for the continent Europe.

```
(life_exp_summ_eur <- data_gapminder %>%
  dplyr::filter(continent == "Europe") %>%
  dplyr::summarise(average = mean(lifeExp),
                   med = median(lifeExp),
                   std = sd(lifeExp),
                   variance = var(lifeExp),
                   iqr = IQR(lifeExp)
  )
)
```

```
## # A tibble: 1 x 5
##   average  med  std variance  iqr
##   <dbl> <dbl> <dbl>    <dbl> <dbl>
## 1   71.9  72.2  5.43    29.5  5.88
```

Grouped Summaries

Question: What are continent-wise summary statistics?

```
(summ_life_exp <- data_gapminder %>%
  dplyr::group_by(continent) %>%
```

```

dplyr::summarise(average = mean(lifeExp),
                  med = median(lifeExp),
                  std = sd(lifeExp),
                  variance = var(lifeExp),
                  iqr = IQR(lifeExp)
                )
)

```

```


## # A tibble: 5 x 6
##   continent average   med   std variance   iqr
##   <fct>      <dbl> <dbl> <dbl>     <dbl> <dbl>
## 1 Africa      48.9  47.8  9.15     83.7  12.0
## 2 Americas    64.7  67.0  9.35     87.3  13.3
## 3 Asia        60.1  61.8 11.9     141.   18.1
## 4 Europe      71.9  72.2  5.43     29.5   5.88
## 5 Oceania     74.3  73.7  3.80     14.4   6.35

```

```

ggplot(data = summ_life_exp,
        aes(x = continent, y = average)) +
  geom_bar(stat = "identity")

```




```
(summ_year_life_exp <- data_gapminder %>%
  dplyr::group_by(year) %>%
  dplyr::summarise(average = mean(lifeExp),
                    med = median(lifeExp),
                    std = sd(lifeExp),
                    variance = var(lifeExp),
                    iqr = IQR(lifeExp)
  )
)
```

```
## # A tibble: 12 x 6
##   year average med std variance iqr
##   <int> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 1952 49.1 45.1 12.2 149. 20.7
## 2 1957 51.5 48.4 12.2 150. 21.8
## 3 1962 53.6 50.9 12.1 146. 21.8
## 4 1967 55.7 53.8 11.7 137. 21.4
## 5 1972 57.6 56.5 11.4 130. 20.7
```

```
## 6 1977 59.6 59.7 11.2 126. 19.9
## 7 1982 61.5 62.4 10.8 116. 18.0
## 8 1987 63.2 65.8 10.6 111. 16.9
## 9 1992 64.2 67.7 11.2 126. 16.5
## 10 1997 65.0 69.4 11.6 134. 18.5
## 11 2002 65.7 70.8 12.3 151. 19.9
## 12 2007 67.0 71.9 12.1 146. 19.3
```

```
ggplot(data = summ_year_life_exp,
       aes(x = year, y = average)) +
  geom_point()
```



One can also perform grouped mutates and grouped filters.