

Introduction to Data Reading and Tidying

Abhinav Anand

Setup

The following discussion assumes we have downloaded R and RStudio. The packages `readr` and `tidyr` need to be installed prior to running the commands below. They are included in the `tidyverse`

1. For downloading R, visit <https://cran.r-project.org/>
2. For downloading RStudio visit <https://www.rstudio.com/>

Reading and Parsing Data Files

The following discussion assumes that all data files referenced are in the same folder as the R codes.

Reading Plain-Text Files (`.csv`, `.tsv` etc.)

We will be working with the following set of files to illustrate the ideas regarding reading real-life, empirical data files.

```
file_fin_risk <- "FMC_T4_read_file_fin_risk.csv"
file_gdppc <- "FMC_T4_read_file_gdppc.csv"
file_US_corp_spread <- "FMC_T4_read_file_US_corp_spread.csv"
```

While we may read files in formats such as `.xls`, `.xlsx` etc. (“excel files”) in R by using the package `readxl`, it is advised by many writers to convert such

files into plain-text .csv format (comma separated format) and then open them by the readr package functions.

read_csv()

read_csv() reads .csv files. For semicolon separated files, read_csv2() function is used.

```
(fin_risk <- readr::read_csv(file_fin_risk)) #file path
```

```
## Parsed with column specification:
```

```
## cols(
```

```
##   Country = col_character(),
```

```
##   Year = col_integer(),
```

```
##   `Risk Points for Foreign Debt as a % of GDP` = col_double(),
```

```
##   `Risk Points for Exchange Rate Stability` = col_double(),
```

```
##   `Risk Points for Debt Service as a % of XGS` = col_double(),
```

```
##   `Risk Points for Current Account as % of XGS` = col_double(),
```

```
##   `Risk Points for International Liquidity` = col_double(),
```

```
##   `Aggregate Financial Risk` = col_double()
```

```
## )
```

```
## # A tibble: 4,380 x 8
```

	Country	Year	`Risk Points for F~`	`Risk Points for ~`	`Risk Points for ~`
	<chr>	<int>	<dbl>	<dbl>	<dbl>
## 1	Albania	1984	5.33	9	NA
## 2	Albania	1985	6	9	NA
## 3	Albania	1986	6	9	NA
## 4	Albania	1987	6	8.42	NA
## 5	Albania	1988	6.25	8	NA
## 6	Albania	1989	6.54	8	NA
## 7	Albania	1990	6.96	8	NA
## 8	Albania	1991	6.75	6.5	NA
## 9	Albania	1992	4.58	5	NA

```
## 10 Albania 1993 4 5 NA
## # ... with 4,370 more rows, and 3 more variables: `Risk Points for Current
## # Account as % of XGS` <dbl>, `Risk Points for International
## # Liquidity` <dbl>, `Aggregate Financial Risk` <dbl>
```

`read_csv` uses the first row as the column names of data. If however, we know this to not be true (sometimes there are a few lines of metadata at the top of the file) we can instruct `read_csv` to refrain from such behavior.

```
(US_corp_spread <- readr::read_csv(file_US_corp_spread))
```

```
## Warning: Missing column names filled in: 'X2' [2], 'X3' [3], 'X4' [4],
## 'X6' [6], 'X7' [7], 'X8' [8], 'X9' [9]
```

```
## Warning: Duplicated column names deduplicated: 'FRED Graph Observations' =>
## 'FRED Graph Observations_1' [5]
```

```
## # A tibble: 41 x 9
##   `FRED Graph Obse~ X2      X3      X4      `FRED Graph Obs~ X6      X7      X8
##   <chr>             <chr> <chr> <chr> <chr>             <chr> <chr> <chr>
## 1 Federal Reserve ~ <NA> <NA> <NA> Federal Reserve~ <NA> <NA> <NA>
## 2 Link: https://fr~ <NA> <NA> <NA> Link: https://f~ <NA> <NA> <NA>
## 3 Help: https://fr~ <NA> <NA> <NA> Help: https://f~ <NA> <NA> <NA>
## 4 Economic Researc~ <NA> <NA> <NA> Economic Resear~ <NA> <NA> <NA>
## 5 Federal Reserve ~ <NA> <NA> <NA> Federal Reserve~ <NA> <NA> <NA>
## 6 <NA>              <NA> <NA> <NA> <NA>              <NA> <NA> <NA>
## 7 AAA              Moody~ <NA> <NA> BAA              Mood~ <NA> <NA>
## 8 <NA>              <NA> <NA> <NA> <NA>              <NA> <NA> <NA>
## 9 Frequency: Annua~ <NA> <NA> <NA> Frequency: Annu~ <NA> <NA> <NA>
## 10 observation_date AAA    <NA> <NA> observation_date BAA    <NA> <NA>
## # ... with 31 more rows, and 1 more variable: X9 <chr>
```

```
(US_corp_spread_skip <- readr::read_csv(file_US_corp_spread,
                                         skip = 10)
)
```

```
## Warning: Missing column names filled in: 'X3' [3], 'X4' [4], 'X7' [7],
```

```
## 'X8' [8]

## Warning: Duplicated column names deduplicated: 'observation_date' =>
## 'observation_date_1' [5]

## # A tibble: 31 x 9
##   observation_date    AAA X3    X4 observation_date_1    BAA X7    X8
##   <date>            <dbl> <chr> <chr> <date>            <dbl> <chr> <chr>
## 1 1985-01-01        12.1 <NA> <NA> 1985-01-01        13.3 <NA> <NA>
## 2 1986-01-01        10.0 <NA> <NA> 1986-01-01        11.4 <NA> <NA>
## 3 1987-01-01         8.36 <NA> <NA> 1987-01-01         9.72 <NA> <NA>
## 4 1988-01-01         9.88 <NA> <NA> 1988-01-01        11.1 <NA> <NA>
## 5 1989-01-01         9.62 <NA> <NA> 1989-01-01        10.6 <NA> <NA>
## 6 1990-01-01         8.99 <NA> <NA> 1990-01-01         9.94 <NA> <NA>
## 7 1991-01-01         9.04 <NA> <NA> 1991-01-01        10.4 <NA> <NA>
## 8 1992-01-01         8.2  <NA> <NA> 1992-01-01         9.13 <NA> <NA>
## 9 1993-01-01         7.91 <NA> <NA> 1993-01-01         8.67 <NA> <NA>
## 10 1994-01-01        6.92 <NA> <NA> 1994-01-01         7.65 <NA> <NA>
## # ... with 21 more rows, and 1 more variable: `BAA-AAA` <dbl>
```

Notes

1. When the data file does not have column names we can use `col_names = FALSE` to tell `read_csv()` not to treat the first row as headings, and instead label them sequentially from X_1 to X_n .
2. While base R has the classic `read.csv()` function to read `.csv` files, usage of `read_csv()` is encouraged since the latter is said to be around 10 times faster than the former. This is critical when file sizes become large. Additionally, the files are read as tibbles and hence retain their readability, flexibility and reproducibility.
3. Excel files can be read with `readxl()`. Files in formats foreign to R, such as Stata files (`.dta`) can be read using the tidyverse package `haven`.
4. R can also write dataframes into a `.csv` file by use of the command

```
write_csv().
```

Tidying Data

```
(gdppc <- readr::read_csv(file_gdppc)) #which format?
```

```
## Parsed with column specification:
```

```
## cols(
```

```
##   .default = col_character()
```

```
## )
```

```
## See spec(...) for full column specifications.
```

```
## # A tibble: 264 x 61
```

```
##   `Series Name`           `Series Code`  `Country Name`      `Country Code`
```

```
##   <chr>                  <chr>          <chr>              <chr>
```

```
## 1 GDP per capita (curren~ NY.GDP.PCAP.CD Afghanistan AFG
```

```
## 2 GDP per capita (curren~ NY.GDP.PCAP.CD Albania ALB
```

```
## 3 GDP per capita (curren~ NY.GDP.PCAP.CD Algeria DZA
```

```
## 4 GDP per capita (curren~ NY.GDP.PCAP.CD American Samoa ASM
```

```
## 5 GDP per capita (curren~ NY.GDP.PCAP.CD Andorra AND
```

```
## 6 GDP per capita (curren~ NY.GDP.PCAP.CD Angola AGO
```

```
## 7 GDP per capita (curren~ NY.GDP.PCAP.CD Antigua and Barb~ ATG
```

```
## 8 GDP per capita (curren~ NY.GDP.PCAP.CD Arab World ARB
```

```
## 9 GDP per capita (curren~ NY.GDP.PCAP.CD Argentina ARG
```

```
## 10 GDP per capita (curren~ NY.GDP.PCAP.CD Armenia ARM
```

```
## # ... with 254 more rows, and 57 more variables: `1960 [YR1960]` <chr>,
```

```
## # `1961 [YR1961]` <chr>, `1962 [YR1962]` <chr>, `1963 [YR1963]` <chr>,
```

```
## # `1964 [YR1964]` <chr>, `1965 [YR1965]` <chr>, `1966 [YR1966]` <chr>,
```

```
## # `1967 [YR1967]` <chr>, `1968 [YR1968]` <chr>, `1969 [YR1969]` <chr>,
```

```
## # `1970 [YR1970]` <chr>, `1971 [YR1971]` <chr>, `1972 [YR1972]` <chr>,
```

```
## # `1973 [YR1973]` <chr>, `1974 [YR1974]` <chr>, `1975 [YR1975]` <chr>,
```

```
## # `1976 [YR1976]` <chr>, `1977 [YR1977]` <chr>, `1978 [YR1978]` <chr>,
```

```
## # `1979 [YR1979]` <chr>, `1980 [YR1980]` <chr>, `1981 [YR1981]` <chr>,
## # `1982 [YR1982]` <chr>, `1983 [YR1983]` <chr>, `1984 [YR1984]` <chr>,
## # `1985 [YR1985]` <chr>, `1986 [YR1986]` <chr>, `1987 [YR1987]` <chr>,
## # `1988 [YR1988]` <chr>, `1989 [YR1989]` <chr>, `1990 [YR1990]` <chr>,
## # `1991 [YR1991]` <chr>, `1992 [YR1992]` <chr>, `1993 [YR1993]` <chr>,
## # `1994 [YR1994]` <chr>, `1995 [YR1995]` <chr>, `1996 [YR1996]` <chr>,
## # `1997 [YR1997]` <chr>, `1998 [YR1998]` <chr>, `1999 [YR1999]` <chr>,
## # `2000 [YR2000]` <chr>, `2001 [YR2001]` <chr>, `2002 [YR2002]` <chr>,
## # `2003 [YR2003]` <chr>, `2004 [YR2004]` <chr>, `2005 [YR2005]` <chr>,
## # `2006 [YR2006]` <chr>, `2007 [YR2007]` <chr>, `2008 [YR2008]` <chr>,
## # `2009 [YR2009]` <chr>, `2010 [YR2010]` <chr>, `2011 [YR2011]` <chr>,
## # `2012 [YR2012]` <chr>, `2013 [YR2013]` <chr>, `2014 [YR2014]` <chr>,
## # `2015 [YR2015]` <chr>, `2016 [YR2016]` <chr>
```

```
head(fin_risk) #which format?
```

```
## # A tibble: 6 x 8
##   Country Year `Risk Points for F~ `Risk Points for E~ `Risk Points for ~
##   <chr>   <int>          <dbl>          <dbl>          <dbl>
## 1 Albania 1984          5.33           9             NA
## 2 Albania 1985           6             9             NA
## 3 Albania 1986           6             9             NA
## 4 Albania 1987           6            8.42          NA
## 5 Albania 1988          6.25           8             NA
## 6 Albania 1989          6.54           8             NA
## # ... with 3 more variables: `Risk Points for Current Account as % of
## #   XGS` <dbl>, `Risk Points for International Liquidity` <dbl>,
## #   `Aggregate Financial Risk` <dbl>
```

The Tidy Format

The tidy format has three characteristics:

1. Each variable is a column

2. Each observation is a row
3. Each value is a cell

`fin_risk` is a tidy dataset, `gdppc` is not.

Not all formats of data are equally good for analysis. For the tidyverse, the best format to work with is the “tidy” format. `dplyr`, `ggplot2` and all the other packages in the tidyverse are designed to work best with tidy data.

```
(fin_risk_tidy <- fin_risk %>%
  dplyr::rename(risk_foreign =
    `Risk Points for Foreign Debt as a % of GDP`) %>%
  dplyr::rename(risk_exchange =
    `Risk Points for Exchange Rate Stability`) %>%
  dplyr::rename(risk_debt =
    `Risk Points for Debt Service as a % of XGS`) %>%
  dplyr::rename(risk_CA =
    `Risk Points for Current Account as % of XGS`) %>%
  dplyr::rename(risk_liq =
    `Risk Points for International Liquidity`) %>%
  dplyr::rename(risk_agg_fin = `Aggregate Financial Risk`)
)
```

```
## # A tibble: 4,380 x 8
##   Country Year risk_foreign risk_exchange risk_debt risk_CA risk_liq
##   <chr>   <int>      <dbl>         <dbl>      <dbl>  <dbl>
## 1 Albania 1984        5.33          9         NA      NA      NA
## 2 Albania 1985         6          9         NA      NA      NA
## 3 Albania 1986         6          9         NA      NA      NA
## 4 Albania 1987         6         8.42        NA      NA      NA
## 5 Albania 1988        6.25          8         NA      NA      NA
## 6 Albania 1989        6.54          8         NA      NA      NA
## 7 Albania 1990        6.96          8         NA      NA      NA
## 8 Albania 1991        6.75         6.5         NA      NA      NA
## 9 Albania 1992        4.58          5         NA      NA      NA
```

```
## 10 Albania 1993 4 5 NA NA NA
## # ... with 4,370 more rows, and 1 more variable: risk_agg_fin <dbl>
```


```
(fin_risk_tidy_summ <- fin_risk_tidy %>%
  dplyr::group_by(Year) %>%
  dplyr::summarise(risk_agg_min =
    min(risk_agg_fin, na.rm = T),
    risk_agg_max =
    max(risk_agg_fin, na.rm = T),
    risk_agg_med =
    median(risk_agg_fin, na.rm = T),
    risk_agg_mean =
    mean(risk_agg_fin, na.rm = T),
    risk_agg_std =
    sd(risk_agg_fin, na.rm = T),
    risk_agg_iqr =
    IQR(risk_agg_fin, na.rm = T)
  )
)
```

```
## # A tibble: 30 x 7
##   Year risk_agg_min risk_agg_max risk_agg_med risk_agg_mean risk_agg_std
##   <int>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>
## 1  1984          0       45.1       22.5       19.3       14.3
## 2  1985          0      116.       26.1       23.8       14.9
## 3  1986          0       44.8       25.9       23.7       12.2
## 4  1987          0       45.5       25.9       23.6       12.3
## 5  1988          0       46.6       26.7       24.1       12.5
## 6  1989          0       46       26.4       24.2       12.5
## 7  1990          0       45.8       26.6       24.8       12.3
## 8  1991          0       45.8       27.7       25.2       12.8
## 9  1992          0       46.0       29.7       25.9       13.2
## 10 1993          0       44.4       29.9       26.0       12.9
```



```
## # ... with 20 more rows, and 1 more variable: risk_agg_iqr <dbl>
```

```
ggplot(data = filter(fin_risk_tidy,  
  Country %in% c("United Kingdom",  
    "United States",  
    "China",  
    "India")  
),  
  mapping = aes(x = Year,  
    y = risk_agg_fin,  
    color = Country)) +  
  geom_point() +  
  geom_line(mapping = aes(linetype = Country))
```



Gathering

This is used to “gather” data from the wide format, to the long format.

```
gdppc %>% head(.)
```

```
## # A tibble: 6 x 61
##   `Series Name`          `Series Code` `Country Name` `Country Code`
##   <chr>                <chr>        <chr>         <chr>
## 1 GDP per capita (current US$) NY.GDP.PCAP.~ Afghanistan   AFG
## 2 GDP per capita (current US$) NY.GDP.PCAP.~ Albania       ALB
## 3 GDP per capita (current US$) NY.GDP.PCAP.~ Algeria       DZA
## 4 GDP per capita (current US$) NY.GDP.PCAP.~ American Samoa ASM
## 5 GDP per capita (current US$) NY.GDP.PCAP.~ Andorra       AND
## 6 GDP per capita (current US$) NY.GDP.PCAP.~ Angola        AGO
## # ... with 57 more variables: `1960 [YR1960]` <chr>, `1961
## #   [YR1961]` <chr>, `1962 [YR1962]` <chr>, `1963 [YR1963]` <chr>, `1964
## #   [YR1964]` <chr>, `1965 [YR1965]` <chr>, `1966 [YR1966]` <chr>, `1967
## #   [YR1967]` <chr>, `1968 [YR1968]` <chr>, `1969 [YR1969]` <chr>, `1970
## #   [YR1970]` <chr>, `1971 [YR1971]` <chr>, `1972 [YR1972]` <chr>, `1973
## #   [YR1973]` <chr>, `1974 [YR1974]` <chr>, `1975 [YR1975]` <chr>, `1976
## #   [YR1976]` <chr>, `1977 [YR1977]` <chr>, `1978 [YR1978]` <chr>, `1979
## #   [YR1979]` <chr>, `1980 [YR1980]` <chr>, `1981 [YR1981]` <chr>, `1982
## #   [YR1982]` <chr>, `1983 [YR1983]` <chr>, `1984 [YR1984]` <chr>, `1985
## #   [YR1985]` <chr>, `1986 [YR1986]` <chr>, `1987 [YR1987]` <chr>, `1988
## #   [YR1988]` <chr>, `1989 [YR1989]` <chr>, `1990 [YR1990]` <chr>, `1991
## #   [YR1991]` <chr>, `1992 [YR1992]` <chr>, `1993 [YR1993]` <chr>, `1994
## #   [YR1994]` <chr>, `1995 [YR1995]` <chr>, `1996 [YR1996]` <chr>, `1997
## #   [YR1997]` <chr>, `1998 [YR1998]` <chr>, `1999 [YR1999]` <chr>, `2000
## #   [YR2000]` <chr>, `2001 [YR2001]` <chr>, `2002 [YR2002]` <chr>, `2003
## #   [YR2003]` <chr>, `2004 [YR2004]` <chr>, `2005 [YR2005]` <chr>, `2006
## #   [YR2006]` <chr>, `2007 [YR2007]` <chr>, `2008 [YR2008]` <chr>, `2009
## #   [YR2009]` <chr>, `2010 [YR2010]` <chr>, `2011 [YR2011]` <chr>, `2012
## #   [YR2012]` <chr>, `2013 [YR2013]` <chr>, `2014 [YR2014]` <chr>, `2015
## #   [YR2015]` <chr>, `2016 [YR2016]` <chr>
```

```

col_yr_1 <- "1960 [YR1960]"
col_yr_end <- "2016 [YR2016]"

(gdppc_tidy <- gdppc %>%
  dplyr::select(-c(`Series Name`,
                    `Series Code`,
                    `Country Code`
                  )
               ) %>%
  dplyr::rename(Country = `Country Name`) %>%
  tidyr::gather(col_yr_1:col_yr_end,
                key = "Year",
                value = "GDP_per_capita"
               ) %>%
  dplyr::arrange(Country)
)

```

```

## # A tibble: 15,048 x 3
##   Country      Year      GDP_per_capita
##   <chr>        <chr>        <chr>
## 1 Afghanistan 1960 [YR1960] 59.7773265084
## 2 Afghanistan 1961 [YR1961] 59.8781528089
## 3 Afghanistan 1962 [YR1962] 58.4928738323
## 4 Afghanistan 1963 [YR1963] 78.7827580363
## 5 Afghanistan 1964 [YR1964] 82.2084438594
## 6 Afghanistan 1965 [YR1965] 101.2904712742
## 7 Afghanistan 1966 [YR1966] 137.899361897
## 8 Afghanistan 1967 [YR1967] 161.3220000885
## 9 Afghanistan 1968 [YR1968] 129.5066538443
## 10 Afghanistan 1969 [YR1969] 129.7985414084
## # ... with 15,038 more rows

```

Spreading

Joining